

# Q & A for the Final Exam

by Baikjin Jung

Gary Geunbae Lee  
CSED342 - Artificial Intelligence

June, 2022

1. **(20 points)** Consider the following grid world game:

This is a stationary MDP with an infinite horizon. The agent can only be in one of the six locations. It gets the reward (or punishment) written in a particular cell when it leaves the cell. It gets a reward of 10 for leaving the bottom-middle square and a reward of  $-100$  (punishment of 100) for leaving the top-left square. In each iteration of the game, the agent has to choose a direction to move. The agent can choose to move either up, down, left, or right. There is a 0.8 probability that it will move in that direction and a 0.1 probability that it will move in either of the neighboring directions. For example, if the agent wants to move up, there is a 0.8 probability that it will move up, a 0.1 probability that it will move left, and a 0.1 probability that it will move right. If the agent bumps into a wall, it stays in its current location and does not get any reward (or punishment).

-100		
	+10	

Figure 1: The given grid world.

- (a) **(8 points)** Perform one step of value iteration and show the resulting value function for each state.

<b>-10</b>	<b>0</b>	<b>0</b>
<b>0</b>	<b>+10</b>	<b>0</b>

Figure 2: **The right answer.**

No partial Credit.

- (b) **(12 points)** What is the value in the top-left state after performing another step of value iteration?

**We take the maximum value of the four possible actions:**

$$\left\{ \begin{array}{ll} \text{left:} & 0.8 \times (0 + (-10)) + 0.1 \times (-100 + 0) + 0.1 \times (-100 + 0) = -28 \\ \text{up:} & 0.8 \times (0 + (-10)) + 0.1 \times (-100 + 0) + 0.1 \times (-100 + 0) = -28 \\ \text{down:} & 0.8 \times (-100 + 0) + 0.1 \times (-100 + 0) + 0.1 \times (0 + (-10)) = -91 \\ \text{right:} & 0.8 \times (-100 + 0) + 0.1 \times (-100 + 0) + 0.1 \times (0 + (-10)) = -91 \end{array} \right.$$

**Therefore, the answer is  $-19$ .**

Partial Credit: 3 points for each direction.



2. **(20 points)** Provided that one is faced with a task to classify a given image of a handwritten digit into one of ten classes representing integer values from 0 to 9, the go-to approach of these days would be using a convolutional neural network (CNN). Let's say that we are given a  $5 \times 5$  pixel image of number 3, for example (the aquamarine pixel is what we are specially interested in.).

Our CNN consists of two layers in the following order:

- a convolution layer with  $2 \times 2$  filter, which is followed by a vectorization procedure (just changing the form),
- and a fully-connected layer that maps a 4-dimensional column vector to a 10-dimensional column vector.

The figure below depicts the above-mentioned layers. Other details are as follows:

- in the original image  $x$ , a colored pixel has a value of 1, and a white pixel has a value of 0;
- the true label  $y = [0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]^T$ ;



- the predicted label  $\hat{y} = [0 \ 0 \ 0 \ 0.5 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]^T$ ;
- all the weight parameters ( $W$ ) of the fully-connected layer have a value of 1.
- the output of the convolution layer

$$\forall n, m \in [1, 4] : c[n, m] = \sum_{u=1}^2 \sum_{v=1}^2 x[n + (u - 1), m + (v - 1)] \cdot k[u, v];$$

- the output of the fully-connected layer  $\hat{y} = \sigma(W \times c)$ , where

$$\sigma(z) = \frac{1}{1 + e^{-z}};$$

- the loss

$$L = \frac{1}{2} \sum_{i=1}^{10} (\hat{y}(i) - y(i))^2.$$

(a) **(10 points)** Calculate  $\frac{\partial L}{\partial c}$ . Refer to the derivatives below if needed.

$$\frac{d}{dz}[f(g(z))] = f'(g(z)) \cdot g'(z) \quad (1)$$

$$\frac{d}{dz} \left[ \frac{1}{x} \right] = -\frac{1}{x^2} \quad (2)$$

$$\frac{d}{dz}[e^z] = e^z \quad (3)$$

$$\begin{aligned} \frac{\partial L}{\partial c} &= \sum_{i=1}^{10} \frac{\partial L}{\partial \hat{y}(i)} \cdot \frac{\partial \hat{y}(i)}{\partial c} \\ &= \left( \sum_{i=1}^{10} (\hat{y}(i) - y(i)) \cdot \hat{y}(i)(1 - \hat{y}(i)) \right) \\ &= -0.125 \end{aligned}$$

Partial Credit:

- Differentiation of the mean squared error: 5 points,
- Differentiation of the sigmoid function: 5 points.

(b) **(10 points)** Calculate the back-propagated gradient for  $k[1, 2]$  when the upper left corner of the filter is on  $x[1, 1]$ , which corresponds to the aquamarine pixel. For differentiation, refer to the derivatives above.

$$\begin{aligned} \nabla k[1, 2] &= \frac{\partial L}{\partial k[1, 2]} \\ &= \sum_{i=1}^{10} \frac{\partial L}{\partial \hat{y}(i)} \cdot \frac{\partial \hat{y}(i)}{\partial k[1, 2]} \\ &= \left( \sum_{i=1}^{10} (\hat{y}(i) - y(i)) \cdot \frac{\partial \hat{y}(i)}{\partial c} \right) \cdot \frac{\partial c}{\partial k[1, 2]} \\ &= \left( \sum_{i=1}^{10} (\hat{y}(i) - y(i)) \cdot \hat{y}(i)(1 - \hat{y}(i)) \right) \cdot W_{i, \cdot} \cdot \frac{\partial c}{\partial k[1, 2]} \\ &= \left( \sum_{i=1}^{10} (\hat{y}(i) - y(i)) \hat{y}(i)(1 - \hat{y}(i)) \right) \cdot x[1, 2] \\ &= -0.125 \end{aligned}$$

No partial Credit.

3. **(20 points)** Use the  $k$ -means algorithm and Euclidean distance to cluster eight data points into  $k = 3$  clusters. The distance matrix based on the Euclidean distance is given in the table below. The coordinates of the data points are:

$$\begin{aligned} x^{(1)} &= (2, 8), & x^{(2)} &= (2, 5), & x^{(3)} &= (1, 2), & x^{(4)} &= (5, 8), \\ x^{(5)} &= (7, 3), & x^{(6)} &= (6, 4), & x^{(7)} &= (8, 4), & x^{(8)} &= (4, 7). \end{aligned}$$

	$x^{(1)}$	$x^{(2)}$	$x^{(3)}$	$x^{(4)}$	$x^{(5)}$	$x^{(6)}$	$x^{(7)}$	$x^{(8)}$
$x^{(1)}$	0	3	6.0828	3	7.0711	5.6569	7.2111	2.2361
$x^{(2)}$	3	0	3.1623	4.2426	5.3852	4.1231	6.0828	2.8284
$x^{(3)}$	6.0828	3.1623	0	7.2111	6.0828	5.3852	7.2801	5.8310
$x^{(4)}$	3	4.2426	7.2111	0	5.3852	4.1231	5	1.4142
$x^{(5)}$	7.0711	5.3852	6.0828	5.3852	0	1.4142	1.4142	5
$x^{(6)}$	5.6569	4.1231	5.3852	4.1231	1.4142	0	2	3.6056
$x^{(7)}$	7.2111	6.0828	7.2801	5	1.4142	2	0	5
$x^{(8)}$	2.2361	2.8284	5.8310	1.4142	5	3.6056	5	0

- (a) **(2 points)** Suppose that we initialize the centroids with  $k$  randomly chosen data points. Let's assume that those points are  $\mu^{(1)} \leftarrow x^{(3)}$ ,  $\mu^{(2)} \leftarrow x^{(4)}$ , and  $\mu^{(3)} \leftarrow x^{(6)}$ . Perform one iteration and assign each data point to the closest cluster.

$$\begin{aligned} c^{(1)} &\leftarrow C^{(2)} \\ c^{(2)} &\leftarrow C^{(1)} \\ c^{(3)} &\leftarrow C^{(1)} \\ c^{(4)} &\leftarrow C^{(2)} \\ c^{(5)} &\leftarrow C^{(3)} \\ c^{(6)} &\leftarrow C^{(3)} \\ c^{(7)} &\leftarrow C^{(3)} \\ c^{(8)} &\leftarrow C^{(2)} \end{aligned}$$

No partial Credit.

- (b) **(4 points)** Next, move the centroids.

$$\begin{aligned} \mu^{(1)} &\leftarrow \frac{1}{2}(x^{(2)} + x^{(3)}) = (1.5, 3.5) \\ \mu^{(2)} &\leftarrow \frac{1}{3}(x^{(1)} + x^{(4)} + x^{(8)}) = (3.67, 7.67) \\ \mu^{(3)} &\leftarrow \frac{1}{3}(x^{(5)} + x^{(6)} + x^{(7)}) = (7, 3.67) \end{aligned}$$

No partial Credit.

- (c) **(7 points)** Calculate the loss function before the first iteration

$$J_0(c^{(1)}, \dots, c^{(m)}, \mu^{(1)}, \dots, \mu^{(k)}) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu^{(c^{(i)})}\|_2^2,$$

where  $c^{(i)}$  is the cluster assigned to the  $i^{\text{th}}$  data point by the algorithm, and the added quantity on the right-hand side is the square of the (given) Euclidean distance between two data points.

$$\begin{aligned} J_0 &= \frac{1}{8} (3^2 + 3.1623^2 + 1.4142^2 + 2^2 + 1.4142^2) \\ &= 3.375 \end{aligned}$$

No partial Credit.


- (d) **(7 points)** Calculate the loss function after the first iteration

$$J_1(c^{(1)}, \dots, c^{(m)}, \mu^{(1)}, \dots, \mu^{(k)}) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu^{(c^{(i)})}\|_2^2.$$

$$\begin{aligned} J_1 &= \frac{1}{8} (2.9 + 2.5 + 2.5 + 1.9 + 0.44 + 1.11 + 1.11 + 0.56) \\ &= 1.625 \end{aligned}$$

No partial Credit.

#### Q4) Most Likely Estimates in HMMs



The Viterbi algorithm finds the most probable sequence of hidden states  $X_{1:T}$ , given a sequence of observations  $e_{1:T}$ . Throughout this question you may assume there are no ties. Recall that for the canonical HMM structure, the Viterbi algorithm performs the following dynamic programming computations:

$$m_t[x_t] = P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

\* Note about dynamic programming: dynamic programming is essentially a recursive relation in which the current value is defined as a function of previously computed values. In this case, the value at time  $t$  is defined as a function of the values at time  $t - 1$ .

(a) [5 points]

For the HMM structure above, which of the following probabilities are maximized by the sequence of states returned by the Viterbi algorithm? Pick **all** correct option(s).


- 1)  $P(X_{1:T})$                       2)  $P(X_T|e_T)$                       3)  $P(X_{1:T}|e_{1:T})$
- 4)  $P(X_{1:T}, e_{1:T})$             5)  $P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1})$     6)  $P(X_1) \prod_{t=2}^T P(X_t|X_{t-1})$
- 7) None of the above

**Answer: 3, 4, 5**

The sequence of states returned by the Viterbi Algorithm maximizes the conditional  $= P(X_{1:T}|e_{1:T})$ . Since,  $P(e_{1:T})$  is just a constant,  $P(X_{1:T}, e_{1:T}) \propto P(X_{1:T}|e_{1:T})$ . Hence, the full joint  $P(X_{1:T}, e_{1:T})$  is also maximized. The third option  $P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1}) = P(X_{1:T}, e_{1:t})$  due to the conditional independences implied by the HMM structure; therefore, this is also maximized.

(b) 5 points each

Consider an HMM structure like the one in part (a) above. Say for all time steps  $t$ , the state  $X_t$  can take on one of the three values  $\{A, B, C\}$ . Then, we can represent the state transitions through the following directed graph, also called a Trellis Diagram.



We wish to formulate the most probable sequence of hidden state query as a graph search problem. Note in the diagram above, dummy nodes S and G have been added to represent the start state and the goal state respectively. Further, the transition from the starting node S to the first state  $X_1$  occurs at time step  $t = 0$ ; transition from  $X_T$  (the last HMM state) to the goal state G occurs at time step  $t = T$ .

**Definition:** Let  $w_{Y \rightarrow Z}^t$ , be the cost of the edge for the transition from state Y at time  $t$  to state Z at time  $t+1$ . For example,  $w_{A \rightarrow B}^1$  is the cost of the edge for transition from state A at time 1 to state B at time 2.

(i) For which **one** of the following values for the weights  $w_{Y \rightarrow Z}^t$ ,  $1 \leq t < T$ , would the minimum cost path be exactly the same as most likely sequence of states computed by the Viterbi algorithm?

- (a)  $w_{Y \rightarrow Z}^t = -P(X_{t+1} = Z | X_t = Y)$       (b)  $w_{Y \rightarrow Z}^t = -P(e_{t+1} | X_{t+1} = Z)P(X_{t+1} = Z | X_t = Y)$   
(c)  $w_{Y \rightarrow Z}^t = -\log(P(X_{t+1} = Z | X_t = Y))$       (d)  $w_{Y \rightarrow Z}^t = -\log(P(e_{t+1} | X_{t+1} = Z)P(X_{t+1} = Z | X_t = Y))$   
(e)  $w_{Y \rightarrow Z}^t = \frac{1}{P(X_{t+1} = Z | X_t = Y)}$       (f)  $w_{Y \rightarrow Z}^t = \frac{1}{P(e_{t+1} | X_{t+1} = Z)P(X_{t+1} = Z | X_t = Y)}$

**Answer: d**



We want the solution to maximize the joint  $P(X_{1:T}, e_{1:T}) = P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1})$ . Since, a search algorithm **minimizes** the cost, we want to pose this problem as a minimization problem. Hence, we can equivalently say that we want to minimize  $\frac{1}{P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1})}$ . A search algorithm in its native form can only work with **additive** costs. Therefore, to turn the above products of probabilities into sums, we take the log.

Hence, we wish to minimize :

$$\log \left( \frac{1}{P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1})} \right) = -\log \left( P(X_1)P(e_1|X_1) \prod_{t=2}^T P(e_t|X_t)P(X_t|X_{t-1}) \right)$$

$$= -\log (P(X_1)P(e_1|X_1)) - \sum_{t=2}^T \log (P(e_t|X_t)P(X_t|X_{t-1})).$$

If for  $t > 1$ , the edge cost is set to  $-\log P(e_{t+1}|X_{t+1} = Z)P(X_{t+1} = Z|X_t = Y)$ , we get the above as the total cost of the path.

(ii) The initial probability distribution of the state at time  $t = 1$  is given  $P(X_1 = Y)$ ,  $Y \in \{A, B, C\}$ . Which **one** of the following should be the value of  $w_{S \rightarrow Y}^0$ ,  $Y \in \{A, B, C\}$  - these are the cost on the edges connecting  $S$  to the states at time  $t = 1$ ?

- |  |  |
|--|--|
| (a) $w_{S \rightarrow Y}^0 = -P(X_1 = Y)$          | (b) $w_{S \rightarrow Y}^0 = -P(e_1 X_1 = Y)P(X_1 = Y)$          |
| (c) $w_{S \rightarrow Y}^0 = -\log(P(X_1 = Y))$    | (d) $w_{S \rightarrow Y}^0 = -\log(P(e_1 X_1 = Y)P(X_1 = Y))$    |
| (e) $w_{S \rightarrow Y}^0 = \frac{1}{P(X_1 = Y)}$ | (f) $w_{S \rightarrow Y}^0 = \frac{1}{P(e_1 X_1 = Y)P(X_1 = Y)}$ |

Answer: d

(iii) Which **one** of the following should be the value of  $w_{Y \rightarrow G}^T$ ,  $Y \in \{A, B, C\}$  – these are the cost on the edges connecting the states at the last time step  $t = T$  to the goal state  $G$ ?

(a)  $w_{Y \rightarrow G}^T = -P(X_T = Y)$

(b)  $w_{Y \rightarrow G}^T = -P(e_T | X_T = Y)P(X_T = Y)$

(c)  $w_{Y \rightarrow G}^T = -\log(P(X_T = Y))$

(d)  $w_{Y \rightarrow G}^T = -\log(P(e_T | X_T = Y)(P(X_T = Y)))$

(e)  $w_{Y \rightarrow G}^T = \frac{1}{P(X_T = Y)}$

(f)  $w_{Y \rightarrow G}^T = \frac{1}{P(e_T | X_T = Y)P(X_T = Y)}$

(g)  $w_{Y \rightarrow G}^T = \alpha$ ,  $\alpha \in \mathbb{R}$  : (some constant)


Answer: g

As long as the costs on all the three edges :  $w_{A \rightarrow G}^T, w_{B \rightarrow G}^T, w_{C \rightarrow G}^T$  is the same, we will get the optimal answer. Note, it does not matter if the constant  $\alpha$  is positive or negative because the total probability is only scaled by a positive constant  $= e^\alpha$  (that is, the total cost of the path is (sum of all log probabilities) +  $w_{Y \rightarrow G}^T (= \alpha)$ ). When you exponentiate this to get the probabilities, you get  $e^{\text{sum of log probabilities}} * e^\alpha$  — which is just a scaling by some positive number).

**Q5) ML: Short Question and Answer Problems [20 points]**

**(a) Parameter Estimation and Smoothing**

For the Bayes' net drawn on the left, A can take on values +a, -a, and B can take values +b and -b. We are given samples (on the right), and we want to use them to estimate  $P(A)$  and  $P(B|A)$ .



**(i) [5 points]** Compute the maximum likelihood estimates for  $P(A)$  and  $P(B|A)$ , and fill in the 2 tables below ((1) ~ (6)).

A	P(A)
+a	(1)
-a	(2)

A	B	P(B A)
+a	+b	(3)
+a	-b	(4)
-a	+b	(5)
-a	-b	(6)

Answer:

(1): 1/10    (2): 9/10

(3): 1/1    (4): 0/1    (5): 4/9    (6): 5/9

**(ii) [5 points]** Compute the estimates for  $P(A)$  and  $P(B|A)$  using Laplace smoothing with strength  $k=2$ , and fill in the 2 tables below ((1) ~ (6)).

<b>A</b>	<b>P(A)</b>
+a	(1)
-a	(2)

<b>A</b>	<b>B</b>	<b>P(B A)</b>
+a	+b	(3)
+a	-b	(4)
-a	+b	(5)
-a	-b	(6)


**Answer:**

(1):  $\frac{3}{14}$    (2):  $\frac{11}{14}$

(3):  $\frac{3}{5}$    (4):  $\frac{2}{5}$    (5):  $\frac{6}{13}$    (6):  $\frac{7}{13}$

**(b) [5 points] Linear Separability**

You are given samples from 2 classes ( . And + ) with each sample being described by 2 features  $f_1$  and  $f_2$ . These samples are plotted in the following figure. You observe that these samples are not *linearly* separable using just these 2 features. Pick the minimal set of features (choices) below that you could use alongside  $f_1$  and  $f_2$ , to *linearly* separate samples from the 2 classes.



Choice A)  $f_1 < 1.5$

Choice B)  $f_1 > -1.5$

Choice C)  $f_2 < -1$

Choice D)  $f_1 > 1.5$

Choice E)  $f_2 < 1$

Choice F)  $f_2 > -1$

Choice G)  $f_1 < -1.5$

Choice H)  $f_2 > 1$

Choice I)  $f_1^2$

Choice J)  $f_2^2$

Choice K)  $|f_1 + f_2|$

Choice L) Even using all these features alongside  $f_1$  and  $f_2$  will not make the samples linearly separable.

**Answer: Choice I and Choice J**

**(c) [5 points] Perceptrons**

In this question, you will perform perceptron updates. You have 2 classes, +1 and -1, and 3 features  $f_0$ ,  $f_1$ ,  $f_2$  for each training point. The +1 class is predicted if  $w \cdot f > 0$  and the -1 class is predicted otherwise.

You start with the weight vector  $w = [1 \ 0 \ 0]$ . In the table below, do a perceptron update for each of the given samples. If the  $w$  vector does not change, write “No Change”, otherwise write down the new  $w$  vector for (i) ~ (iii).

<b>f0</b>	<b>f1</b>	<b>f2</b>	<b>Class</b>	<b>Updated w</b>
1	7	8	-1	(i)
1	6	8	-1	(ii)
1	9	6	+1	(iii)

**Answer:**

(i):  $[0 \ -7 \ -8]$

(ii): Unchanged (no misclassification)

(iii):  $[1 \ 2 \ -2]$