

Lecture 16

Discrete Probability

Suha Kwak

suha.kwak@postech.ac.kr

Dept. of Computer Science and Engineering

POSTECH

An Introduction to Discrete Probability

Probability of an Event

- Key terms

- An **experiment** is a procedure that yields one of a given set of possible outcomes.
- The **sample space** of the experiment is the set of possible outcomes.
- An **event** is a subset of the sample space.

- Example

- Random experiment
 - Tossing a single six-sided dice
- Sample space
 - $S = \{\square \cdot \square \cdot \cdot \square \square \square \square\}$
- Event
 - A subset of the sample space
 - “The number on the face > 3 ” = $\{\square \square \square \square\}$
 - S = Sure event
 - ϕ = Impossible event

Probability of an Event

- Definition (*by Pierre-Simon Laplace*)
 - If S is a finite sample space of equally likely outcomes, and E is an event (*i.e.*, a subset of S) then the probability of E is

$$P(E) = \frac{|E|}{|S|}.$$

- For every event E , we have $0 \leq P(E) \leq 1$.
 - This follows directly from the definition because $0 \leq P(E) = |E|/|S| \leq |S|/|S| \leq 1$, since $0 \leq |E| \leq |S|$.
- Example 1
 - An urn contains four blue balls and five red balls.
 - What is the probability that a ball chosen from the urn is blue?
 - Answer: $\frac{4}{9}$

Probability of an Event

- Example 2
 - There are many lotteries that award prizes to people who correctly choose a set of six numbers out of the first n positive integers, where n is usually between 30 and 60.
 - What is the probability that a person picks the correct six numbers out of 40?
 - Answer:
 - The number of ways to choose six numbers out of 40 is $C(40,6) = 40!/(34! 6!) = 3,838,380$.
 - Hence the probability of picking a winning combination is $\frac{1}{3,838,380} \approx 0.00000026$.

Probability of an Event

- Example 3
 - What is the probability that the numbers 11, 4, 17, 39, and 23 are drawn in that order from a bin with 50 balls labeled with the numbers 1, 2, ..., 50 if
 - A. The ball selected is not returned to the bin.
 - B. The ball selected is returned to the bin before the next ball is selected.
 - Answer:
 - Use the product rule in each case.
 - A. (*Sampling without replacement*) The probability is $\frac{1}{254,251,200}$ since there are $50 \cdot 49 \cdot 48 \cdot 47 \cdot 46 = 254,251,200$ ways to choose the five balls.
 - B. (*Sampling with replacement*) Since $50^5 = 312,500,000$ the probability is $\frac{1}{50^5} = \frac{1}{312,500,000}$.

Complements and Unions of Events

- Theorem: Probability of the complementary event
 - Let E be an event in sample space S .
 - The probability of the event $\bar{E} = S - E$, the **complementary** event of E , is given by $P(\bar{E}) = 1 - P(E)$.
 - Proof
 - Using the fact that $|\bar{E}| = |S| - |E|$,
- Example
 - A sequence of 10 bits is chosen randomly.
 - What is the probability that at least one of these bits is 0?

$$P(E) = 1 - P(\bar{E}) = 1 - \frac{|\bar{E}|}{|S|} = 1 - \frac{1}{2^{10}} = \frac{1023}{1024}$$

Complements and Unions of Events

- Theorem: Probability of a union of events
 - Let E_1 and E_2 be events in the sample space S , then

$$P(E_1 \cup E_2) = P(E_1) + P(E_2) - P(E_1 \cap E_2)$$

- Proof
 - Given the inclusion-exclusion formula,
 $|A \cup B| = |A| + |B| - |A \cap B|.$
 - It follows that

$$\begin{aligned} P(E_1 \cup E_2) &= \frac{|E_1 \cup E_2|}{|S|} = \frac{|E_1| + |E_2| - |E_1 \cap E_2|}{|S|} \\ &= \frac{|E_1|}{|S|} + \frac{|E_2|}{|S|} - \frac{|E_1 \cap E_2|}{|S|} \\ &= P(E_1) + P(E_2) - P(E_1 \cap E_2) \end{aligned}$$

Complements and Unions of Events

- Example
 - Suppose we select an integer at random from the set of positive integers not exceeding 100.
 - What is the probability that the selected integer is divisible by either 2 or 5?
 - Solution:
 - Let E_1 be the event that the integer is divisible by 2 and E_2 be the event that it is divisible by 5.
 - Then the event that the integer is divisible by 2 or 5 is $E_1 \cup E_2$ and $E_1 \cap E_2$ is the event that it is divisible by both 2 and 5.
 - It follows that

$$\begin{aligned}P(E_1 \cup E_2) &= P(E_1) + P(E_2) - P(E_1 \cap E_2) \\&= \frac{50}{100} + \frac{20}{100} - \frac{10}{100} = \frac{3}{5}\end{aligned}$$

Probability Theory

Summary

- Assigning Probabilities
- Conditional Probability
- Independence
- Bernoulli Trials and the Binomial Distribution
- Random Variables

Assigning Probabilities

- Laplace's definition of probability
 - It assumes that all outcomes are equally likely.
 - A more general definition is required to avoid the restriction.
- Assigning probabilities
 - Let S be a sample space of an experiment with a finite number of outcomes.
 - Assign a probability $P(s)$ to each outcome s , so that:
 - $0 \leq P(s) \leq 1$ for each $s \in S$
 - $\sum_{s \in S} P(s) = 1$
 - The function P from the set of all outcomes of the sample space S is called a **probability distribution**.

Assigning Probabilities

- Example
 - What probabilities should we assign to the outcomes H (heads) and T (tails) when a fair coin is flipped?
 - What probabilities should be assigned to these outcomes when the coin is biased so that heads comes up twice as often as tails?
 - Solution:
 - For a fair coin, the probabilities for H and T are both $\frac{1}{2}$.
 - For the biased coin: $P(H) = 2 \cdot P(T)$, and since $P(H) + P(T) = 1$, $P(H) = \frac{2}{3}$ and $P(T) = \frac{1}{3}$.

Uniform Distribution

- Definition
 - Suppose that S is a set with n elements.
 - The **uniform distribution** assigns the probability $\frac{1}{n}$ to each element of S .
 - Note that we could have used Laplace's definition here.
- Example
 - For a fair coin: $P(H) = P(T) = \frac{1}{2}$
 - For a fair six-faced dice: $P(1) = P(2) = \dots = P(6) = \frac{1}{6}$

Probability of an Event

- The probability of the event E is the sum of the probabilities of the outcomes in E :

$$P(E) = \sum_{s \in E} P(s)$$

- Now no assumption is being made about the distribution.
- Example
 - Suppose that a die is biased so that  appears twice as often as each other, but the other five faces appear equally.
 - What is the probability that an odd number appears when we roll this die?

$$P(\text{double dot}) = \frac{2}{7}, \text{ and } P(\text{one dot}) = P(\text{two dots}) = P(\text{three dots}) = P(\text{four dots}) = P(\text{five dots}) = \frac{1}{7}$$

$$P(E) = P(\text{one dot or double dot or four dots}) = P(\text{one dot}) + P(\text{double dot}) + P(\text{four dots}) = \frac{4}{7}$$

Combinations of Events

- Theorem
 - If E_1, E_2, \dots is a sequence of pairwise *disjoint* events in a sample space S , then

$$P\left(\bigcup_i E_i\right) = \sum_i P(E_i)$$

Conditional Probability

- Definition
 - Let E and F be events with $P(F) > 0$. The conditional probability of E given F , denoted by $P(E|F)$, is defined as:
- Example
 - A bit string of length 4 is sampled from a uniform distribution.
 - What is the probability that it contains at least two consecutive 0s, *given* that its first bit is a 0?
 - Let E be the event that the bit string contains at least two consecutive 0s, and F be the event that the first bit is a 0.
 - $E \cap F = \{0000, 0001, 0010, 0011, 0100\}$, so $P(E \cap F) = 5/16$.
 - $P(F) = 1/2$ since there are 8 bit strings starting with a 0.
 - Hence, $P(E|F) = 5/16 \cdot 2 = 5/8$

Independence

- Definition
 - The events E and F are independent if and only if
$$P(E \cap F) = P(E) \cdot P(F)$$
- Example
 - Suppose E is the event that a randomly generated bit string of length 4 begins with a 1 and F is the event that this bit string contains an even number of 1s.
 - Are E and F independent if the 16 bit strings of length 4 are equally likely?
 - They are independent; why?
 - There are eight bit strings of length 4 beginning with a 1, and eight bit strings of length 4 that contain an even number of 1s.
 - Since the number of bit strings of length 4 is 16, $P(E) = P(F) = \frac{1}{2}$
 - Since $E \cap F = \{1111, 1100, 1010, 1001\}$, $P(E \cap F) = \frac{1}{4}$.

Pairwise and Mutual Independence

- Definition (*Pairwise independence*)
 - The events E_1, E_2, \dots, E_n are pairwise independent if and only if

$$P(E_i \cap E_j) = P(E_i) \cdot P(E_j)$$

for all pairs i and j with $1 \leq i < j \leq n$.

- Definition (*Mutual independence*)
 - The events E_1, E_2, \dots, E_n are mutually independent if and only if

$$P(E_{i_1} \cap E_{i_2} \cap \dots \cap E_{i_m}) = P(E_{i_1}) \cdot P(E_{i_2}) \cdots P(E_{i_m})$$

whenever i_j ($j = 1, 2, \dots, m$) are integers with

$$1 \leq i_1 < i_2 < \dots < i_m \leq n \text{ and } m \geq 2$$

Pairwise and Mutual Independence

- Example

- A fair 6-sided dice is thrown twice.
- Three events of our interest:
 - E_1 : The first roll is a 3.
 - E_2 : The second roll is 4.
 - E_3 : The sum of the two rolls is 7.
- The three events are pairwise independent since
 - $P(E_1) = P(E_2) = P(E_3) = \frac{1}{6}$
 - $P(E_1 \cap E_2) = P(E_2 \cap E_3) = P(E_1 \cap E_3) = \frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$
- However, the events are not mutually independent since
 - $P(E_1 \cap E_2 \cap E_3) = \frac{1}{36} \neq \left(\frac{1}{6}\right)^3$

Bernoulli Trials

- Definition
 - Suppose an experiment can have only two possible outcomes.
 - *E.g.*, the flipping of a coin or the random generation of a bit.
 - Each performance of the experiment is called a **Bernoulli trial**.
 - One outcome is called a *success* and the other a *failure*.
 - If p is the probability of success and q the probability of failure, then $p + q = 1$.

Many problems involve determining the probability of k successes when an experiment consists of n mutually independent Bernoulli trials.

Bernoulli Trials

- Example
 - A coin is biased so that the probability of heads is $2/3$.
 - What is the probability that exactly 4 heads occur when the coin is flipped 7 times?
 - Solution:
 - There are $2^7 = 128$ possible outcomes.
 - The number of ways 4 of the 7 flips can be heads is $C(7,4)$.
 - The probability of each of the outcomes is $\left(\frac{2}{3}\right)^4 \left(\frac{1}{3}\right)^3$ since the 7 flips are independent.
 - Hence, the probability that exactly 4 heads occur is

$$C(7,4) \cdot \left(\frac{2}{3}\right)^4 \left(\frac{1}{3}\right)^3 = \frac{35 \cdot 16}{3^7} = \frac{560}{2187}$$

Bernoulli Trials

- Theorem
 - The probability of exactly k successes in n independent Bernoulli trials, with probability of success p and probability of failure $q = 1 - p$, is
$$C(n, k) \cdot p^k \cdot q^{n-k}$$
- Proof
 - The outcome of n Bernoulli trials is an n -tuple (t_1, t_2, \dots, t_n) , where each t_i is either S (success) or F (failure).
 - The probability of each outcome of n trials consisting of k successes and $n - k$ failures (in any order) is $p^k q^{n-k}$.
 - Because there are a $C(n, k)$ number of n -tuple of S s and F s that contain exactly k S s, the probability of k successes is $C(n, k) \cdot p^k \cdot q^{n-k}$.
 - We denote this probability by $b(k: n, p) = C(n, k) \cdot p^k \cdot q^{n-k}$, a function of k given n and p .

Random Variables

- Definition

- A **random variable** is a function from the sample space of an experiment to the set of real numbers.
- That is, a random variable assigns a real number to each possible outcome.

A random variable is a function.
It is not a variable, and it is not random!

In the late 1940s W. Feller and J.L. Doob flipped a coin to see whether both would use “random variable” or the more fitting “chance variable.” Unfortunately, Feller won and the term “random variable” has been used ever since...

Random Variables

- Definition
 - The distribution of a random variable X on a sample space S is the set of pairs $(r, P(X = r))$ for all $r \in X(S)$, where $P(X = r)$ is the probability that X takes the value r .
- Example: A coin flipped 3 times
 - Let $X(t)$ be the random variable that equals the number of heads that appear when t is the outcome.
 - Then $X(t)$ takes on the following values:
 - $X(HHH) = 3, X(TTT) = 0$
 - $X(HHT) = X(HTH) = X(THH) = 2$
 - $X(TTH) = X(THT) = X(HTT) = 1$
 - Each of the 8 possible outcomes has probability $1/8$.
 - So, the distribution of $X(t)$ is $P(X = 3) = 1/8$, $P(X = 2) = 3/8$, $P(X = 1) = 3/8$, and $P(X = 0) = 1/8$.

Bayes' Theorem

Bayes' Theorem

- Bayes' theorem
 - Suppose that E and F are events from a sample space S such that $P(E) \neq 0$ and $P(F) \neq 0$. Then:

$$P(F|E) = \frac{P(E|F)P(F)}{P(E|F)P(F) + P(E|\bar{F})P(\bar{F})}$$

- Proof

$$\begin{aligned} P(F|E) &= \frac{P(E \cap F)}{P(E)} = \frac{P(E \cap F)}{P(E \cap F) + P(E \cap \bar{F})} \\ &= \frac{P(E|F)P(F)}{P(E|F)P(F) + P(E|\bar{F})P(\bar{F})} \end{aligned}$$

Bayes' Theorem

- Example
 - We have two boxes.
 - The first box contains 2 green balls and 7 red balls.
 - The second contains 4 green balls and 3 red balls.
 - You select one of the boxes at random, then select a ball from the selected box at random.
 - If you have a red ball, what is the probability that you selected a ball from the first box.
- Solution
 - E : The event that you have chosen a red ball
 - F : The event that you have chosen the first box

$$\begin{aligned} P(F|E) &= \frac{P(E|F)P(F)}{P(E|F)P(F) + P(E|\bar{F})P(\bar{F})} \\ &= \frac{(7/9)(1/2)}{(7/9)(1/2) + (3/7)(1/2)} = \frac{49}{76} \end{aligned}$$

Generalized Bayes' Theorem

- Generalized Bayes' Theorem
 - Suppose that E is an event from a sample space S and that F_1, F_2, \dots, F_n are *mutually exclusive* events such that

$$\bigcup_{i=1}^n F_i = S$$

- Assume that $P(F_i) \neq 0$ for $i = 1, 2, \dots, n$.
- Then

$$P(F_j | E) = \frac{P(E | F_j)P(F_j)}{\sum_{i=1}^n P(E | F_i)P(F_i)}$$

A Little Taste of Machine Learning

- Interpreting Bayes' theorem

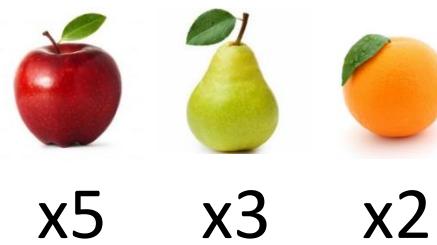
$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

- The event of our interest A
 - Typically latent (*i.e.*, hidden)
- The event as an observation B
- Prior probability $P(A)$
 - Probability of the event A with no observation
 - Based only on our *prior* knowledge about A
- Likelihood $P(B|A)$
 - Probability of observing B when A happens
- Posterior probability $P(A|B)$
 - Probability of A if we observed B

posterior \propto *likelihood* \times *prior*

A Little Taste of Machine Learning

- Interpreting Bayes' theorem (cont'd)



x5 x3 x2

[prior]

$$P(\text{apple}) = 0.5, \quad P(\text{pear}) = 0.3, \quad P(\text{orange}) = 0.2$$

[likelihood]

$$P(\text{tongue} | \text{apple}) = 0.3, \quad P(\text{tongue} | \text{pear}) = 0.5, \quad P(\text{tongue} | \text{orange}) = 0.9$$

[posterior \propto likelihood \times prior]

$$P(\text{apple} | \text{tongue}) \propto 0.15, \quad P(\text{pear} | \text{tongue}) \propto 0.15, \quad P(\text{orange} | \text{tongue}) \propto 0.18$$

A Little Taste of Machine Learning

- Bayesian spam filtering
 - Suppose that we have an initial set B of spam messages and set G of non-spam messages.
 - Datasets like B and G are called *training data* in the context of machine learning
 - We can use this information along with Bayes' theorem to predict the probability that a new email message is spam.
 - We look at a particular word w , and count the number of times that it occurs in B and in G : $n_B(w)$ and $n_G(w)$.
 - Let S be the event that the message is spam, and E be the event that the message contains the word w .
 - Estimated probability that a spam message contains w :
$$P(E|S) = n_B(w)/|B|$$
 - Estimated probability that a non-spam message contains w :
$$P(E|\bar{S}) = n_G(w)/|G|$$

A Little Taste of Machine Learning

- Bayesian spam filtering (cont'd)
 - For spam filtering, we compute the posterior probability $P(S|E)$ for the message and categorize it as spam if the probability is larger than a pre-defined threshold.
 - $P(S|E)$ indicates the probability that the message is spam when it contains the word w .

$$\begin{aligned} P(S|E) &= \frac{P(E|S)P(S)}{P(E)} = \frac{P(E|S)P(S)}{P(E|S)P(S) + P(E|\bar{S})P(\bar{S})} \\ &= \frac{P(E|S)}{P(E|S) + P(E|\bar{S})} \quad \text{Here we assuming that} \\ &\quad \text{the prior probabilities are uniform.} \end{aligned}$$

If we have data on the frequency of spam messages,
we can obtain a better estimate of the priors.

A Little Taste of Machine Learning

- Bayesian spam filtering: A specific example
 - We find that the word “Rolex” occurs in 250 out of 2000 spam messages and occurs in 5 out of 1000 non-spam messages.
 - Estimate the probability that an incoming message with “Rolex” is spam.
 - Suppose our threshold for rejecting the email is 0.9.
 - Is the message a spam according to our Bayesian spam filter?

$$P(\text{"Rolex"}|S) = \frac{250}{2000} = 0.125, \quad P(\text{"Rolex"}|\bar{S}) = \frac{5}{1000} = 0.005$$

$$P(S|\text{"Rolex"}) = \frac{P(\text{"Rolex"}|S)}{P(\text{"Rolex"}|S) + P(\text{"Rolex"}|\bar{S})} = \frac{0.125}{0.125 + 0.005} \approx 0.962$$

Since the probability $P(S|\text{"Rolex"})$ is larger than the threshold, our spam filtering system will categorize the message as spam and reject it!

A Little Taste of Machine Learning

- Bayesian spam filtering with using multiple words
 - Accuracy can be improved by considering more than one word as evidence.
 - Consider the case where E_1 and E_2 denote the events that the message contains the words w_1 and w_2 respectively.
 - For simplicity, we make the assumption that $P(S) = 0.5$ (the uniform prior) and E_1 and E_2 are *conditionally independent* given S :

$$P(E_1 \cap E_2 | S) = P(E_1 | S)P(E_2 | S)$$

- Then the posterior probability of S is given by

$$P(S | E_1 \cap E_2) = \frac{P(E_1 | S)P(E_2 | S)}{P(E_1 | S)P(E_2 | S) + P(E_1 | \bar{S})P(E_2 | \bar{S})}$$

A Little Taste of Machine Learning

- Bayesian spam filtering with using multiple words:
A specific example
 - We have 2000 spam messages and 1000 non-spam messages.
 - The word “stock” occurs 400 times in the spam messages and 60 times in the non-spam.
 - The word “undervalued” occurs in 200 spam messages and 25 non-spam.

$$P(E_1|S) = \frac{400}{2000} = 0.2, \quad P(E_1|\bar{S}) = \frac{60}{1000} = 0.06$$

$$P(E_2|S) = \frac{200}{2000} = 0.1, \quad P(E_2|\bar{S}) = \frac{25}{1000} = 0.025$$

$$P(S|E_1 \cap E_2) = \frac{(0.2)(0.1)}{(0.2)(0.1) + (0.06)(0.025)} \approx 0.93$$

A Little Taste of Machine Learning

- Bayesian spam filtering in general
 - The more words we consider, the more accurate the spam filter we can obtain.
 - With the independence assumption if we consider k words:

$$P(S | \cap_{i=1}^k E_i) = \frac{\prod_{i=1}^k P(E_i | S)}{\prod_{i=1}^k P(E_i | S) + \prod_{i=1}^k P(E_i | \bar{S})}$$

- We can further improve the filter by considering pairs of words as a single block of certain types of strings.

Expected Value and Variance

Expected Value

- Definition
 - The **expected value** (or **expectation** or **mean**) of the random variable $X(s)$ on the sample space S is equal to

$$E(X) = \sum_{s \in S} P(s) \cdot X(s)$$

- Example
 - Let X be the number that comes up when a fair die is rolled.
 - What is the expected value of X ?
 - Solution
 - The random variable X takes the values 1, 2, 3, 4, 5, or 6.
 - Each has probability 1/6.
 - It follows that

$$E(X) = \frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 2 + \dots + \frac{1}{6} \cdot 6 = \frac{21}{6} = \frac{7}{2}$$

Expected Value

- Theorem 2
 - Suppose that n mutually independent Bernoulli trials are performed, where p is the probability of success on each trial.
 - Then the expected number of successes in this case is np .
 - Proof
 - Let X be the random variable equal to the number of success in n trials.
 - Then,

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

- Hence,

$$\begin{aligned} E(X) &= \sum_{k=1}^n k \cdot P(X = k) \\ &= \sum_{k=1}^n k \cdot \binom{n}{k} p^k (1 - p)^{n-k} \end{aligned}$$

Expected Value

- Theorem 2
 - Proof (cont'd)

$$\begin{aligned} E(X) &= \sum_{k=1}^n k \cdot P(X = k) = \sum_{k=1}^n k \cdot \binom{n}{k} p^k (1-p)^{n-k} \\ &= \sum_{k=1}^n n \cdot \binom{n-1}{k-1} p^k (1-p)^{n-k} \\ &= np \sum_{k=1}^n \binom{n-1}{k-1} p^{k-1} (1-p)^{n-k} \\ &= np \sum_{j=0}^{n-1} \binom{n-1}{j} p^j (1-p)^{n-1-j} \\ &= np \cdot (p + 1 - p)^{n-1} = np \cdot 1 = np \end{aligned}$$

Shift the index of sum with $j = k - 1$.

Recall the binomial theorem:

$$(x + y)^n = \sum_{j=0}^n \binom{n}{j} x^{n-j} y^j$$

Linearity of Expectations

- Theorem
 - X_i are random variables on S where $i = 1, 2, \dots, n$
 - a and b are real numbers.
 - Then
 - (i) $E(X_1 + X_2 + \dots + X_n) = E(X_1) + E(X_2) + \dots + E(X_n)$
 - (ii) $E(aX + b) = aE(X) + b$
 - Proof of (i)
 - When $n = 2$,
$$\begin{aligned} E(X_1 + X_2) &= \sum_{s \in S} P(s)(X_1(s) + X_2(s)) \\ &= \sum_{s \in S} P(s)X_1(s) + \sum_{s \in S} P(s)X_2(s) = E(X_1) + E(X_2) \end{aligned}$$
 - We can generalize the result for $n > 2$ by math induction.

Linearity of Expectations

- Theorem

- X_i are random variables on S where $i = 1, 2, \dots, n$
- a and b are real numbers.
- Then

$$\begin{aligned}(i) \quad & E(X_1 + X_2 + \cdots + X_n) = E(X_1) + E(X_2) + \cdots + E(X_n) \\(ii) \quad & E(aX + b) = aE(X) + b\end{aligned}$$

- Proof of (ii)

$$\begin{aligned}E(aX + b) &= \sum_{s \in S} P(s)(aX(s) + b) \\&= a \sum_{s \in S} P(s)X(s) + b \sum_{s \in S} P(s) \\&= aE(X) + b \cdot 1 = aE(X) + b\end{aligned}$$

Average-Case Computational Complexity

- The average-case complexity of an algorithm can be found by computing the expected value of a random variable.
 - Let the sample space of an experiment be the set of possible inputs a_j where $j = 1, 2, \dots, n$.
 - Let the random variable X be the assignment to a_j of the number of operations used by the algorithm when given a_j as input.
 - Assign a probability $P(a_j)$ to each possible input value a_j .
 - The expected value of X is the average-case computational complexity of the algorithm.

$$E(X) = \sum_{j=1}^n P(a_j)X(a_j)$$

Average-Case Computational Complexity

- Linear search

```
Procedure linear_search( $x$ : integer,  $\{a_1, a_2, \dots, a_n\}$ : integers)
```

```
     $i := 1$ 
```

```
    while  $i \leq n$  and  $x \neq a_i$ 
```

```
         $i := i + 1$ 
```

```
        if  $i \leq n$  then  $\ell := i$ 
```

```
    else  $\ell := 0$ 
```

```
return  $\ell$ 
```

- Suppose that
 - The probability that x is in the list is p ,
 - It is equally likely that x is any of the n elements of the list.
- Then, what is the average-case computational complexity of this algorithm?

Average-Case Computational Complexity

- Linear search (cont'd)
 - There are $n + 1$ possible types of input:
 - One type for each of the n numbers on the list
 - One additional type for the numbers not on the list.
 - Recall that we need:
 - $2i + 1$ comparisons if x equals the i^{th} element of the list.
 - $2n + 2$ comparisons if x is not on the list.
 - The probability that x equals a_j is p/n and the probability that x is not in the list is $q = 1 - p$.
 - The average-case computational complexity of the linear search algorithm is:

$$\begin{aligned} E &= 3p/n + 5p/n + \cdots + (2n + 1)p/n + (2n + 2)q \\ &= (p/n)(3 + 5 + \cdots + (2n + 1)) + (2n + 2)q \\ &= (p/n)((n + 1)^2 - 1) + (2n + 2)q \\ &= p(n + 2) + (2n + 2)q. \end{aligned}$$

For Independent Random Variables

- Definition

- The random variables X and Y on a sample space S are independent if

$$P(X = r_1 \text{ and } Y = r_2) = P(X = r_1) \cdot P(Y = r_2)$$

- Theorem

- If X and Y are independent variables on a sample space S , then $E(XY) = E(X) \cdot E(Y)$.
 - Proof
 - The event $XY = r$ is the disjoint union of the events $X = r_1$ and $Y = r_2$ over all $r_1 \in X(S)$ and $r_2 \in Y(S)$ with $r = r_1 \cdot r_2$.

$$\begin{aligned} E(XY) &= \sum_{r \in XY(S)} r \cdot P(XY = r) \\ &= \sum_{r_1 \in X(S)} \sum_{r_2 \in Y(S)} r_1 r_2 \cdot P(X = r_1 \text{ and } Y = r_2) \end{aligned}$$

For Independent Random Variables

- Theorem
 - Proof (cont'd)

$$\begin{aligned} E(XY) &= \sum_{r_1 \in X(S)} \sum_{r_2 \in Y(S)} r_1 r_2 \cdot P(X = r_1 \text{ and } Y = r_2) \\ &= \sum_{r_1 \in X(S)} \sum_{r_2 \in Y(S)} r_1 r_2 \cdot P(X = r_1) \cdot P(Y = r_2) \\ &= \sum_{r_1 \in X(S)} \left(r_1 \cdot P(X = r_1) \cdot \sum_{r_2 \in Y(S)} r_2 \cdot P(Y = r_2) \right) \\ &= \sum_{r_1 \in X(S)} r_1 \cdot P(X = r_1) \cdot E(Y) \\ &= E(Y) \left(\sum_{r_1 \in X(S)} r_1 \cdot P(X = r_1) \right) = E(Y)E(X) \end{aligned}$$

Variance

- Deviation
 - The **deviation** of X at $s \in S$ is $X(s) - E(X)$, the difference between the value of X and the mean of X .
- Definition: Variance and Standard deviation
 - Let X be a random variable on the sample space S .
 - The **variance** of X , denoted by $V(X)$ is
$$V(X) = \sum_{s \in S} (X(s) - E(X))^2 \cdot P(s)$$
 - That is $V(X)$ is the weighted average of the square of the deviation of X .
 - The **standard deviation** of X , denoted by $\sigma(X)$ is defined to be
$$\sigma(X) = \sqrt{V(X)}$$

Variance

- Theorem
 - If X is a random variable on a sample space S , then
$$V(X) = E(X^2) - E(X)^2$$
- Corollary
 - If X is a random variable on a sample space S and $E(X) = \mu$,
$$V(X) = E((X - \mu)^2)$$
- Example
 - For a random variable X ,
 - $X(t) = 1$ if a Bernoulli trial t is a success,
 - $X(t) = 0$ if it is a failure.
 - p is the probability of success and q is that of failure.
 - What is the variance of the random variable X ?

$$V(X) = E(X^2) - E(X)^2 = p - p^2 = p(1 - p) = pq$$

